

Sieci neuronowe i sztuczna inteligencja – laboratorium 9

Monika Błyszcz, 236623

Kod:

```
1 # import os utilities
2 import os
3 import requests, zipfile, io
4
5 # import numpy
6 import numpy as np
7
8 # images
9 import skimage
10 from skimage import transform
11 from skimage.color import rgb2gray
12
13 # Import the `pyplot` module
14 import matplotlib.pyplot as plt
15
16 # import random
17 import random
18
19 # import tf
20 import tensorflow as tf
21 import tf_slim as slim
22 import keras
23 tf.compat.v1.disable_eager_execution()
24
25
26
27 # function to load data
28 def load_data(data_directory):
29     """Loads sign images data from their folder.
30
31     Returns:
32         images: list of images, i.e., signs
33         labels: list of labels, i.e., signs IDs
34     """
35     # We need back labels and the row images
36     images = []
37     labels = []
38
39     # We have one folder per sign type
40     directories = []
41     for d in os.listdir(data_directory):
42         if os.path.isdir(os.path.join(data_directory, d)):
43             directories.append(d)
44
45     # In each folder there are not only images but also csv
description
```

```

46     # files
47     for d in directories:
48         label_directory = os.path.join(data_directory, d)
49         file_names = [
50             os.path.join(label_directory, f)
51             for f in os.listdir(label_directory)
52             if f.endswith(".ppm")
53         ]
54
55         for f in file_names:
56             images.append(skimage.io.imread(f))
57             labels.append(int(d))
58
59     return images, labels
60
61
62 ROOT_PATH = os.getcwd()
63
64 # Download training data
65
66 train_data_directory = os.path.join(ROOT_PATH, "Training")
67 test_data_directory = os.path.join(ROOT_PATH, "Testing")
68
69 images, labels = load_data(train_data_directory)
70 test_images, test_labels = load_data(test_data_directory)
71
72 print(labels)
73
74 ## The following commented lines were reported in the
75 DataCamp materials
76 ## but they does not work here
77 # print(images.ndim)
78 # print(images.size)
79 images[0]
79 print(len(images))
80 print(len(labels))
81
82 # this should be a bar plot but an histogram with the same
83 number of
84 # bins that that unique levels of the labels list should
85 be fine :-)
86
87 unique_labels = set(labels)
88 n_labels = max(unique_labels) + 1
89
90 # Make a histogram with 62 bins of the `labels` data
91 plt.hist(labels, n_labels)

```

```

89
90 # Show the plot
91 plt.show()
92
93 # Determine the (random) indexes of the images that you
   want to see
94 traffic_signs = [300, 2250, 3650, 4000]
95
96 # Fill out the subplots with the random images that you
   defined
97 for i in range(len(traffic_signs)):
98     plt.subplot(1, 4, i + 1)
99     plt.axis('off')
100    plt.imshow(images[traffic_signs[i]])
101    plt.subplots_adjust(wspace=0.5)
102    plt.show()
103    print(
104        "shape: {0}, min value: {1}, max value: {2}".
        format(
105            images[traffic_signs[i]].shape,
106            images[traffic_signs[i]].min(),
107            images[traffic_signs[i]].max()
108        )
109    )
110
111
112 ##-----
113
114 # Plot a grid with a sample of all the signs
115 plt.figure(figsize=(15, 15))
116
117 i = 1
118
119 for label in unique_labels:
120     # pick the first image for the label.
121     #
122     # The index() method searches an element in the list
   and returns its
123     # index. In simple terms, index() method finds the
   given element in
124     # a list and returns its position. However, if the
   same element is
125     # present more than once, index() method returns its
   smallest/first
126     # position.
127     image = images[labels.index(label)]

```

```

128
129     # We have 62 images. Hence, define a 64 grid sub-
    plots
130     plt.subplot(8, 8, i)
131
132     # Don't include axes
133     plt.axis('off')
134
135     # Add a title to each subplot
136     #
137     # The count() method returns the number of elements
    with the
138     # specified value.
139     plt.title("Label {0} ({1})".format(label, labels.
    count(label)))
140
141     # Add 1 to the counter
142     i += 1
143
144     # Plot this first image
145     plt.imshow(image)
146
147 plt.show()
148
149 # To tackle the differing image sizes, you're going to
    rescale the images
150 images_28 = [
151     transform.resize(image, (28, 28))
152     for image in images
153 ]
154
155 # Convert `images28` to an array
156 images_28 = np.array(images_28)
157
158 # Convert `images28` to grayscale
159 images_28 = rgb2gray(images_28)
160
161 for i in range(len(traffic_signs)):
162     plt.subplot(1, 4, i + 1)
163     plt.axis('off')
164     plt.imshow(images_28[traffic_signs[i]], cmap="gray")
165     plt.subplots_adjust(wspace=0.5)
166
167 plt.show()
168
169 # Test set

```

```

170 # Transform the images to 28 by 28 pixels
171 test_images_28 = [
172     transform.resize(image, (28, 28))
173     for image in test_images
174 ]
175 # Convert to grayscale
176 test_images_28 = rgb2gray(np.array(test_images_28))
177
178
179 # Lets start tensorflow!!
180
181 # Define placeholders for the inputs and labels
182 x = tf.compat.v1.placeholder(dtype=tf.float32, shape=[
183     None, 28, 28])
184 y = tf.compat.v1.placeholder(dtype=tf.int32, shape=[None
185 ])
186
187 # Flatten the images for the inputs of ANN
188 images_flat = tf.keras.layers.Flatten()(x)
189
190 # Fully connected layer output is 62 as the different
191 # signs
192 # this will be the network!!
193 logits = slim.layers.fully_connected(images_flat, 62, tf.
194     nn.relu)
195
196 # Define a loss function
197 loss = tf.reduce_mean(
198     tf.nn.sparse_softmax_cross_entropy_with_logits(
199         labels=y,
200         logits=logits
201     )
202 )
203
204 # Neural Network
205 #
206 # Define an optimizer
207 train_op = tf.compat.v1.train.AdamOptimizer(learning_rate
208     =0.001).minimize(loss)
209
210
211 # Convert logits to label indexes.
212 # NOTE: this will be the final classifier which output
213 # will be the
214 # predicted labels!!
215 correct_pred = tf.argmax(logits, 1)
216

```



```

210 # Define an accuracy metric
211 accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.
    float32))
212
213 print("images_flat: ", images_flat)
214 print("logits: ", logits)
215 print("loss: ", loss)
216 print("predicted_labels: ", correct_pred)
217
218 # run the Graph
219 tf.random.set_seed(1234)
220
221 with tf.compat.v1.Session() as sess:
222     # initialize all the variables
223     sess.run(tf.compat.v1.global_variables_initializer())
224     losses = []
225     error_train = []
226     error_test = []
227
228     # epoch
229     for i in range(201):
230
231         # run the optimizer, accordingly to the loss
defined, feeding
232         # the actual graph with the input we want. In
this case all the
233         # samples every time.
234
235         # NOTE: this update the weights every time, i.e.
the logits,
236         # i.e. the correct_pred!!!
237         _, loss_value = sess.run(
238             [train_op, loss],
239             feed_dict={x: images_28, y: labels}
240         )
241
242         # Just print the loss every 10 epoch
243         losses.append(loss_value)
244         if i % 10 == 0:
245             print("Loss: ", loss_value)
246
247         # Run predictions against the full train set.
248         predicted_train = sess.run(
249             [correct_pred],
250             feed_dict={x: images_28}
251         )[0]

```

```

252         # Calculate mean test error
253         train_error = 1 - np.mean([
254             int(y == y_)
255             for y, y_ in zip(labels, predicted_train)
256         ])
257         error_train.append(train_error)
258
259
260         # Run predictions against the full test set.
261         predicted_test = sess.run(
262             [correct_pred],
263             feed_dict={x: test_images_28}
264         )[0]
265         # Calculate mean test error
266         test_error = 1 - np.mean([
267             int(y == y_)
268             for y, y_ in zip(test_labels, predicted_test)
269         ])
270         error_test.append(test_error)
271
272         # NOTE: if de-indented the session will be closed and
273         so you cannot
274         #         run the sess.run() call
275
276         # Pick 10 random images
277         sample_indexes = random.sample(range(len(images_28
278         )), 10)
279         sample_images = [images_28[i] for i in sample_indexes
280         ]
281         sample_labels = [labels[i] for i in sample_indexes]
282
283         # To have predictions we have to run the "
284         correct_pred" operation
285         # inside the session, feeding the sample we would
286         like to predict
287         predicted = sess.run([correct_pred], feed_dict={x:
288         sample_images})[0]
289
290         # Print the real and predicted labels
291         print(sample_labels)
292         print(predicted)
293
294         # Display the predictions and the ground truth
295         visually.
296         fig = plt.figure(figsize=(10, 10))
297         for i in range(len(sample_images)):

```

```

291
292     # i starts from 0!!
293     truth = sample_labels[i]
294     prediction = predicted[i]
295     color = 'green' if truth == prediction else 'red'
296
297     plt.subplot(5, 2, 1 + i)
298     plt.axis('off')
299
300     plt.text(
301         x=40, y=10,
302         s="Truth:          {0}\nPrediction: {1}".format
303     (
304         truth, prediction
305     ),
306     fontsize=12,
307     color=color
308 )
309     plt.imshow(sample_images[i], cmap="gray")
310
311 plt.show()
312
313 # Print the accuracy
314 print("Final test error: {:.3f}".format(test_error))
315
316 plt.plot(error_train, "b", error_test, "r--")
317 plt.axvline(
318     x=error_test.index(min(error_test)),
319     color="g", linestyle='--'
320 )
321 plt.ylabel('Overall classification error')
322 plt.xlabel("Epochs")
323 plt.title("Training (blue) and test (red) errors by
epoch")
324 plt.show()
325
326
327 #Predicted
328 predicted1 = sess.run([correct_pred], feed_dict={x:
test_images_28})[0]
329
330 #Calculate correct matches
331 match_count = sum([int (y == y_) for y, y_ in zip(
test_labels,predicted1)])
332

```


Zad 1

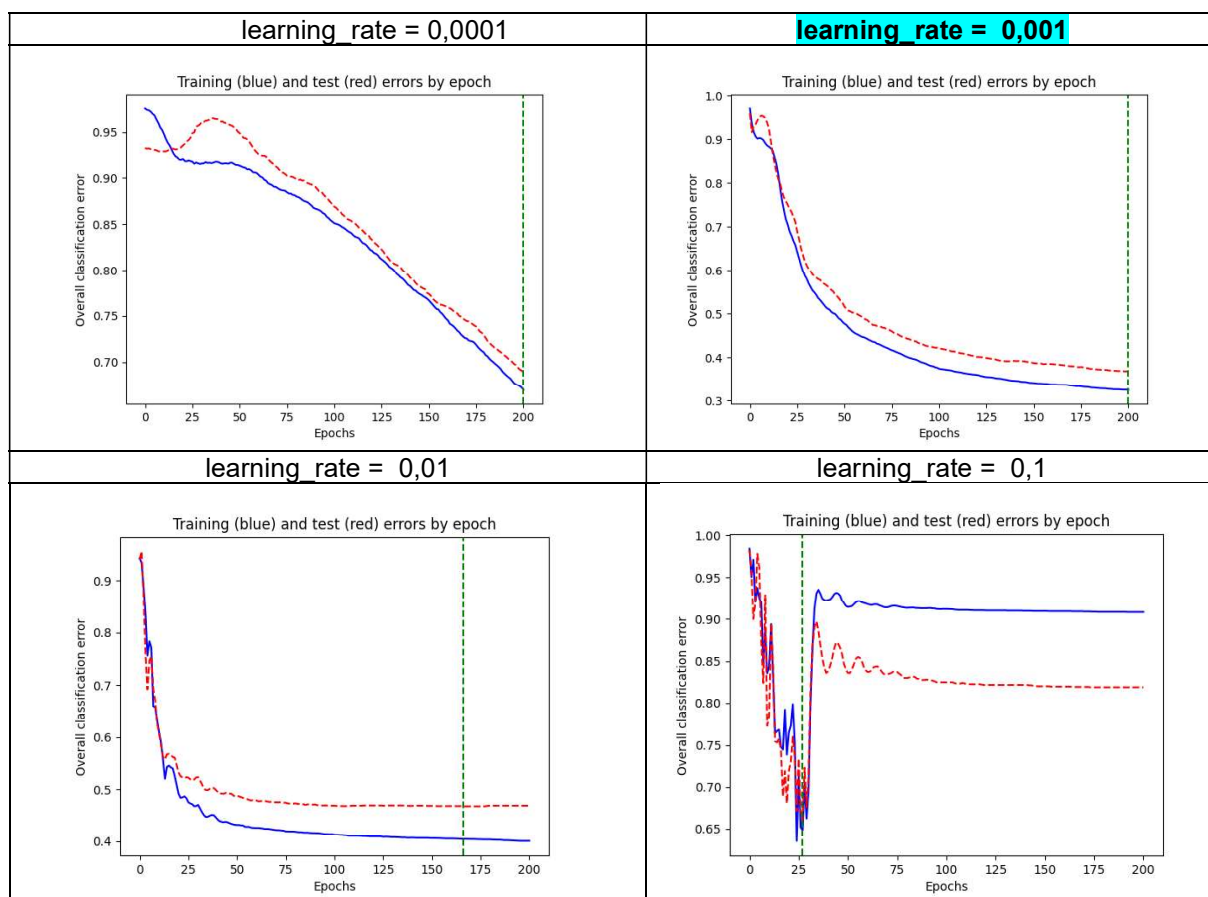
[illegible]

Zad 2

```
200 # Neural Network
201 #
202 # Define an optimizer
203 train_op = tf.compat.v1.train.AdamOptimizer(learning_rate=0.001).minimize(loss)
204
```

learning_rate = 0,000001

learning_rate = 0,00001



Najlepsze wyniki dopasowania uzyskano dla learning rate = 0,001. Wykresy treningowe (Training) oraz testowe (Test) niemal się pokrywają. Największe różnice wyszły dla learning rate = 0,1. Im wyższa wartość, tym bardziej się rozmiągają wynik